

WEEK 10

Ôn Tập – Final review


Loop

```
loopy:  
    # a in %rdi, n in %esi  
    movl    $0, %ecx  
    movl    $0, %edx  
    testl  %esi, %esi  
    jle    .L3  
.L6:  
    movslq %edx,%rax  
    movl    (%rdi,%rax,4), %eax  
    cmpl    %eax, %ecx  
    cmovl   %eax, %ecx  
    addl    $1, %edx  
    cmpl    %ecx, %esi  
    jg     .L6  
.L3:  
    movl    %ecx, %eax  
    ret
```

Fill the blank of corresponding C code.

```
int loopy(int a[], int n)  
{  
    int i;  
    int x = 0;  
    for(i = 0; x < n; i++) {  
        if (x < a[i])  
            x = a[i];  
    }  
    return x;  
}
```

Cache



Parameter	Description
Fundamental parameters	
$S = 2^s$	Number of sets
E	Number of lines per set
$B = 2^b$	Block size (bytes)
$m = \log_2(M)$	Number of physical (main memory) address bits
Derived quantities	
$M = 2^m$	Maximum number of unique memory addresses
$s = \log_2(S)$	Number of <i>set index bits</i>
$b = \log_2(B)$	Number of <i>block offset bits</i>
$t = m - (s + b)$	Number of <i>tag bits</i>
$C = B \times E \times S$	Cache size (bytes), not including overhead such as the valid and tag bits

cache

The following table gives the parameters for a number of different caches. Your task is to fill in the missing fields in the table. Recall that m is the number of physical address bits, C is the cache size (number of data bytes), B is the block size in bytes, E is the associativity, S is the number of cache sets, t is the number of tag bits, s is the number of set index bits, and b is the number of block offset bits.

m	C	B	E	S	t	s	b
32	2048	8	1	<u>256</u>	21	8	3
32	2048	<u>4</u>	<u>4</u>	128	23	7	2
32	1024	2	8	64	<u>25</u>	<u>6</u>	1
32	1024	<u>32</u>	2	16	23	4	<u>5</u>

Cache hit - miss

Suppose we have a system with the following properties:

- The memory is byte addressable.
- Memory accesses are to 1-byte words (not to 4-byte words).
- Addresses are 12 bits wide.
- The cache is two-way set associative ($E = 2$), with a 4-byte block size ($B = 4$) and four sets ($S = 4$).

The contents of the cache are as follows, with all addresses, tags, and values given in hexadecimal notation:

Set index	Tag	Valid	Byte 0	Byte 1	Byte 2	Byte 3
0	00	1	40	41	42	43
	83	1	FE	97	CC	D0
1	00	1	44	45	46	47
	83	0	---	---	---	---
2	00	1	48	49	4A	4B
	40	0	---	---	---	---
3	FF	1	9A	C0	03	FF
	00	0	---	---	---	---

CT	CI	CI	CO	CO							
11	10	9	8	7	6	5	4	3	2	1	0

Operation	Address	Hit	Read (or unknown)
Read	0x834	No	Unknown
Write	0x836	Yes	Unknown
Read	0xFFD	Yes	0xC0

Xác định giá trị

4-way set associative cache																								
Index	Tag	V	Bytes 0–3			Tag	V	Bytes 0–3			Tag	V	Bytes 0–3			Tag	V	Bytes 0–3						
0	F0	1	ED	32	0A	A2	8A	1	BF	80	1D	FC	14	1	EF	09	86	2A	BC	0	25	44	6F	1A
1	BC	0	03	3E	CD	38	A0	0	16	7B	ED	5A	BC	1	8E	4C	DF	18	E4	1	FB	B7	12	02
2	BC	1	54	9E	1E	FA	B6	1	DC	81	B2	14	00	0	B6	1F	7B	44	74	0	10	F5	B8	2E
3	BE	0	2F	7E	3D	A8	C0	1	27	95	A4	74	C4	0	07	11	6B	D8	BC	0	C7	B7	AF	C2
4	7E	1	32	21	1C	2C	8A	1	22	C2	DC	34	BC	1	BA	DD	37	D8	DC	0	E7	A2	39	BA
5	98	0	A9	76	2B	EE	54	0	BC	91	D5	92	98	1	80	BA	9B	F6	BC	1	48	16	81	0A
6	38	0	5D	4D	F7	DA	BC	1	69	C2	8C	74	8A	1	A8	CE	7F	DA	38	1	FA	93	EB	48
7	8A	1	04	2A	32	6A	9E	0	B1	86	56	0E	CC	1	96	30	47	F2	BC	1	F8	1D	42	30

CT	CI	CI	CI	CO	CO	CO							
12	11	10	9	8	7	6	5	4	3	2	1	0	

13bit Addr	Read 0x071A	Block offset (CO)	0x2
		Index (CI)	0x6
		Cache tag (CT)	0x38
		Cache hit? (Y/N)	Y
		Cache byte return	0xEB

Xác định địa chỉ

4-way set associative cache																								
Index	Tag	V	Bytes 0–3			Tag	V	Bytes 0–3			Tag	V	Bytes 0–3			Tag	V	Bytes 0–3						
0	F0	1	ED	32	0A	A2	8A	1	BF	80	1D	FC	14	1	EF	09	86	2A	BC	0	25	44	6F	1A
1	BC	0	03	3E	CD	38	A0	0	16	7B	ED	5A	BC	1	8E	4C	DF	18	E4	1	FB	B7	12	02
2	BC	1	54	9E	1E	FA	B6	1	DC	81	B2	14	00	0	B6	1F	7B	44	74	0	10	F5	B8	2E
3	BE	0	2F	7E	3D	A8	C0	1	27	95	A4	74	C4	0	07	11	6B	D8	BC	0	C7	B7	AF	C2
4	7E	1	32	21	1C	2C	8A	1	22	C2	DC	34	BC	1	BA	DD	37	D8	DC	0	E7	A2	39	BA
5	98	0	A9	76	2B	EE	54	0	BC	91	D5	92	98	1	80	BA	9B	F6	BC	1	48	16	81	0A
6	38	0	5D	4D	F7	DA	BC	1	69	C2	8C	74	8A	1	A8	CE	7F	DA	38	1	FA	93	EB	48
7	8A	1	04	2A	32	6A	9E	0	B1	86	56	0E	CC	1	96	30	47	F2	BC	1	F8	1D	42	30


Read 0x16E8

Block offset (CO)	0x0
Index (CI)	0x2
Cache tag (CT)	0xB7
Cache hit? (Y/N)	N
Cache byte return	---

x86-64/Linux Stack Frame

- Current Stack Frame (“Top” to Bottom)

- “Argument build:”
Parameters for function about to call
- Local variables
If can’t keep in registers
- Saved register context
- Old frame pointer (optional)



- Caller Stack Frame

- Return address
 - Pushed by `call` instruction
- Arguments for this call

Example: incr

```
long incr(long *p, long val) {  
    long x = *p;  
    long y = x + val;  
    *p = y;  
    return x;  
}
```

Register	Use(s)
%rdi	Argument p
%rsi	Argument val , y
%rax	x , Return value


```
incr:  
    movq    (%rdi), %rax  
    addq    %rax, %rsi  
    movq    %rsi, (%rdi)  
    ret
```

Example: Calling `incr` #1


```
long call_incr() {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $15213, 8(%rsp)  
    movl    $3000, %esi  
    leaq    8(%rsp), %rdi  
    call    incr  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

Initial Stack Structure



Resulting Stack Structure




Example: Calling `incr` #2

```
long call_incr() {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $15213, 8(%rsp)  
    movl    $3000, %esi  
    leaq    8(%rsp), %rdi  
    call    incr  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

Stack Structure




Register	Use(s)
%rdi	&v1
%rsi	3000

Example: Calling `incr` #3

```
long call_incr() {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $15213, 8(%rsp)  
    movl    $3000, %esi  
    leaq    8(%rsp), %rdi  
    call    incr  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

Stack Structure




Register	Use(s)
%rdi	&v1
%rsi	3000

Example: Calling `incr` #4

```
long call_incr() {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return v1+v2;  
}
```


```
call_incr:  
    subq    $16, %rsp  
    movq    $15213, 8(%rsp)  
    movl    $3000, %esi  
    leaq    8(%rsp), %rdi  
    call    incr  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

Stack Structure



Register	Use(s)
%rax	Return value

Updated Stack Structure




Example: Calling `incr` #5

```
long call_incr() {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return v1+v2;  
}
```


```
call_incr:  
    subq    $16, %rsp  
    movq    $15213, 8(%rsp)  
    movl    $3000, %esi  
    leaq    8(%rsp), %rdi  
    call    incr  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

Updated Stack Structure



Register	Use(s)
%rax	Return value

Final Stack Structure



Disassembly of last(long u, long v)

u in %rdi, v in %rsi

1	0000000000400540 <last>:								
2	400540: 48 89 f8	mov %rdi,%rax	L1: u						
3	400543: 48 0f af c6	imul %rsi,%rax	L2: u*v						
4	400547: c3	retq	L3: Return						

Disassembly of last(long x)

x in %rdi

5	0000000000400548 <first>:								
6	400548: 48 8d 77 01	lea 0x1(%rdi),%rsi	F1: x+1						
7	40054c: 48 83 ef 01	sub \$0x1,%rdi	F2: x-1						
8	400550: e8 eb ff ff ff	callq 400540 <last>	F3: Call last(x-1,x+1)						
9	400555: f3 c3	repz retq	F4: Return						
.	.								
10	400560: e8 e3 ff ff ff	callq 400548 <first>	M1: Call first(10)						
11	400565: 48 89 c2	mov %rax,%rdx	M2: Resume						


Label	PC	Instruction	State values (at beginning)						Description
			%rdi	%rsi	%rax	%rsp	*%rsp		
M1	0x400560	callq	10	—	—	0x7fffffff820	—	Call first(10)	
F1									
F2									
F3									
L1									
L2									
L3									
F4									
M2									

Starting with the calling of first(10) by main, fill in the following table to trace instruction execution through to the point where the program returns back to main

Label	PC	Instruction	State values (at beginning)						Description
			%rdi	%rsi	%rax	%rsp	*%rsp		
M1	0x400560	callq	10	—	—	0x7fffffff820	—	Call first(10)	
F1	0x400548	lea	10	—	—	0x7fffffff818	0x400565	Entry of first	
F2	0x40054c	sub	10	11	—	0x7fffffff818	0x400565		
F3	0x400550	callq	9	11	—	0x7fffffff818	0x400565	Call last(9, 11)	
L1	0x400540	mov	9	11	—	0x7fffffff810	0x400555	Entry of last	
L2	0x400543	imul	9	11	9	0x7fffffff810	0x400555		
L3	0x400547	retq	9	11	99	0x7fffffff810	0x400555	Return 99 from last	
F4	0x400555	repz repq	9	11	99	0x7fffffff818	0x400565	Return 99 from first	
M2	0x400565	mov	9	11	99	0x7fffffff820	—	Resume main	

Structures & Alignment


- Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



Alignment Principles

- Aligned Data

- Primitive data type requires κ bytes
- Address must be multiple of κ
- Required on some machines; advised on x86-64

- Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory trickier when datum spans 2 pages

- Compiler

- Inserts gaps in structure to ensure correct alignment of fields


Specific Cases of Alignment (x86-64)

- 1 byte: **char**, ...
 - no restrictions on address
- 2 bytes: **short**, ...
 - lowest 1 bit of address must be 0_2
- 4 bytes: **int**, **float**, ...
 - lowest 2 bits of address must be 00_2
- 8 bytes: **double**, **long**, **char ***, ...
 - lowest 3 bits of address must be 000_2
- 16 bytes: **long double** (GCC on Linux)
 - lowest 4 bits of address must be 0000_2

Satisfying Alignment with Structures

- Within structure:
 - Must satisfy each element's alignment requirement
- Overall structure placement
 - Each structure has alignment requirement K
 - $K =$ Largest alignment of any element
 - Initial address & structure length must be multiples of K
- Example:
 - $K = 8$, due to **double** element


```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```



Meeting Overall Alignment Requirement

- For largest alignment requirement K
- Overall structure must be multiple of K


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Struct data alignment

Show how the struct would appear on a 64-bit machine (primitives of size k are k -byte aligned). Label the bytes that belong to the various fields with their names and clearly mark the end of the struct. Use hatch marks or x's to indicate bytes that are allocated in the struct but are not used.

```
struct {
    char a[9]; // 9 bytes
    short b[3]; // 6 bytes
    float c; // 4 bytes
    char d; // 1 byte
    int e; // 4 bytes
    char *f; // 8 bytes (pointer size on 64-bit machine)
    short g; // 2 bytes
} foo;
```

a	x	b	c	d	x	x	x	e	x	x	x	x	f	g	x	x	x	x	x
0 1 2 3 4 5 6 7 8 9	1 1 1 1 1 1 1 1 1	0 1 2 3 4 5 6 7 8 9	1 1 1 1 1 1 1 1 1	2 2 2 2 2 2 2 2 2	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	4 4 4 4 4 4 4 4 4 4	4 4 4 4 4 4 4 4 4 4	4 4 4 4 4 4 4 4 4 4	4 4 4 4 4 4 4 4 4 4	4 4 4 4 4 4 4 4 4 4	


Array Allocation

Basic Principle


$T \mathbf{A}[L];$

- Array of data type T and length L
- Contiguously allocated region of $L * \text{sizeof}(T)$ bytes in memory


`char string[12];`



`int val[5];`



`double a[3];`



`char *p[3];`



Array find M?

```
1 void transpose(long A[M][M]) {  
2     long i, j;  
3     for (i = 0; i < M; i++)  
4         for (j = 0; j < i; j++) {  
5             long t = A[i][j];  
6             A[i][j] = A[j][i];  
7             A[j][i] = t;  
8         }  
9 }
```

```
1 .L6:  
2     movq    (%rdx), %rcx  
3     movq    (%rax), %rsi  
4     movq    %rsi, (%rdx)  
5     movq    %rcx, (%rax)  
6     addq    $8, %rdx  
7     addq    $120, %rax  
8     cmpq    %rdi, %rax  
9     jne     .L6
```

- %rax và %rdx được dùng như con trỏ, ta thấy dòng 6, %rdx tăng lên mỗi lần 8, vậy %rdx chính là A[i][j].
- Thanh ghi còn lại %rax chính là A[j][i]
- Ta thấy %rax tăng mỗi lần là 120, vậy nên M =120/8=15.

Array

Consider the following source code, where M and N are constants declared with `#define`:

```
long P[M][N];
long Q[N][M];

long sum_element(long i, long j) {
    return P[i][j] + Q[j][i];
}
```

Formular for calculating the column sum of Array

$$\&A[i][j] = \sum_{k=0}^{M-1} A[i][k]$$

```
long sum_element(long i, long j)
i in %rdi, j in %rsi
sum_element:
1   leaq    0(%rdi), %rdx
2   subq    %rdi, %rdx
3   addq    %rsi, %rdx
4   leaq    (%rsi,%rsi,4), %rax
5   addq    %rax, %rdi
6   movq    Q(%rdi), %rax
7   addq    P(%rdx), %rax
8
9   ret
```

Use your reverse engineering skills to determine the values of M and N based on this assembly code.

2. $8i$
3. $7i$
4. $7i+j$
5. $5j$
6. $5j+i$
7. $Q+8(5j+i)$
8. $P+8(7i+j)$

N=7
M=5

Array

```
int array1 [H] [J];  
int array2 [J] [H];  
  
int copy_array(int x, int y)  
{  
    array2 [y] [x] = array1 [x] [y];  
    return 1;  
}
```

J=?

H=?

```
# On entry:  
# %edi = x  
# %esi = y  
#  
copy_array:  
    movslq %edi, %rdi  
    movslq %esi, %rsi  
    movq %rdi, %rax  
    leaq (%rsi, %rsi, 2), %rdx  
    salq $5, %rax  
    subq %rdi, %rax  
    leaq (%rdi, %rdx, 2), %rdx  
    addq %rsi, %rax  
    movl array1(%rax, 4), %eax  
    movl %eax, array2(%rdx, 4)  
    movl $1, %eax  
ret
```